

Understanding Black-box Predictions via Influence Functions

June 12, 2021

Content

- ① Introduction
- ② Influence functions
- ③ Scaling up
- ④ Applications
- ⑤ Conclusion

Introduction

Given a high-accuracy, black-box model, and a predict from it. What did the model make this prediction?

1. Make better decisions.
2. Improve the model.
3. Discover new science.
4. Provide end-users explanations.

Introduction

Existing methods:

- Treat model as fixed.
- Explain prediction w.r.t (with respect to) model parameters or test input.

In this paper:

- Treat model as function of training data.
- Explain prediction w.r.t the training data "most responsible" for predict.

To formalize the impact of a training point on a prediction:

- What would happen if we didn't have this training point?
- Or if the values of this training point were changed slightly?

Influence functions

Problem

The symbols:

- \mathcal{X} : input space
- \mathcal{Y} : output space
- $\{z_i = (x_i, y_i)\}_{i=1}^n$: training points in $\mathcal{X} \times \mathcal{Y}$
- For a point z and parameters $\theta \in \Theta$, let $L(z, \theta)$ be the loss, and let $\frac{1}{n} \sum_{i=1}^n L(z_i, \theta)$ be the empirical risk.

The empirical risk minimizer is given by:

$$\hat{\theta} := \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n L(z_i, \theta)$$

Assume: Empirical risk is twice-differentiable and strictly convex in θ

Influence functions

Upweighting a training point

- New parameters: $F(\epsilon) = \hat{\theta}_{\epsilon,z} := \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n L(z_i, \theta) + \epsilon L(z, \theta)$
- Define:

$$\mathcal{I}_{\text{up,params}}(z) := \left. \frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} \right|_{\epsilon=0} = H_{\hat{\theta}}^{-1} \nabla_{\theta} L(z, \hat{\theta})$$

with: $H_{\hat{\theta}} := \frac{1}{n} \sum_{i=1}^n \nabla_{\theta}^2 L(z_i, \hat{\theta})$ is the Hessian and is positive definite (PD) by assumption.

- Linear approximation (Taylor): $F(\epsilon) \approx F(0) + \epsilon * \mathcal{I}_{\text{up,params}}(z)$
- So $\hat{\theta}_{\epsilon,z} - \hat{\theta} \approx \epsilon * \mathcal{I}_{\text{up,params}}(z)$ and $\hat{\theta}_{-z} - \hat{\theta} \approx -\frac{1}{n} \mathcal{I}_{\text{up,params}}(z)$

Influence function

Deriving $\mathcal{I}_{\text{up,params}}(z) \stackrel{\text{def}}{=} \left. \frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} \right|_{\epsilon=0}$

- Recall that $\hat{\theta}$ minimizes the empirical risk:

$$R(\theta) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n L(z_i, \theta) \quad (1)$$

- We further assume that R is twice-differentiable and strongly convex in θ :

$$H_{\hat{\theta}} \stackrel{\text{def}}{=} \nabla^2 R(\hat{\theta}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta}^2 L(z_i, \hat{\theta}) \quad (2)$$

- The perturbed parameters $\hat{\theta}_{\epsilon,z}$ can be written as:

$$\hat{\theta}_{\epsilon,z} = \arg \min_{\theta \in \Theta} \{R(\theta) + \epsilon L(z, \theta)\} \quad (3)$$

Influence functions

Deriving $\mathcal{I}_{\text{up,params}}(z) \stackrel{\text{def}}{=} \left. \frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} \right|_{\epsilon=0}$

- Define: $\Delta_{\epsilon} = \hat{\theta}_{\epsilon,z} - \hat{\theta}$. We have:

$$\frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} = \frac{d\Delta_{\epsilon}}{d\epsilon} \quad (4)$$

- Since $\hat{\theta}_{\epsilon,z}$ is minimizer of (3), let us examine its firstorder optimality conditions:

$$0 = \nabla R(\hat{\theta}_{\epsilon,z}) + \epsilon \nabla L(z, \hat{\theta}_{\epsilon,z}) \quad (5)$$

Influence functions

Deriving $\mathcal{I}_{\text{up,params}}(z) \stackrel{\text{def}}{=} \left. \frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} \right|_{\epsilon=0}$

- Since $\hat{\theta}_{\epsilon,z} \rightarrow \theta$ as $\epsilon \rightarrow 0$, we perform a Taylor expansion:

$$0 \approx [\nabla R(\hat{\theta}) + \epsilon \nabla L(z, \hat{\theta})] + \left[\nabla^2 R(\hat{\theta}) + \epsilon \nabla^2 L(z, \hat{\theta}) \right] \Delta_{\epsilon} \quad (6)$$

- We have $\nabla R(\hat{\theta}) = 0$. Keeping only $O(\epsilon)$ terms, we have:

$$\Delta_{\epsilon} \approx -\nabla^2 R(\hat{\theta})^{-1} \nabla L(z, \hat{\theta}) \epsilon \quad (7)$$

- Finally:

$$\left. \frac{d\hat{\theta}_{\epsilon,z}}{d\epsilon} \right|_{\epsilon=0} = -H_{\hat{\theta}}^{-1} \nabla L(z, \hat{\theta}) \stackrel{\text{def}}{=} \mathcal{I}_{\text{up,params}}(z) \quad (8)$$

Influence functions

Deriving other functions

Influence of upweighting z on the loss at the test point z_{test} :

$$\begin{aligned}\mathcal{I}_{\text{up,loss}}(z, z_{\text{test}}) &\stackrel{\text{def}}{=} \left. \frac{dL(z_{\text{test}}, \hat{\theta}_{\epsilon, z})}{d\epsilon} \right|_{\epsilon=0} \\ &= \nabla_{\theta} L(z_{\text{test}}, \hat{\theta})^{\top} \left. \frac{d\hat{\theta}_{\epsilon, z}}{d\epsilon} \right|_{\epsilon=0} \\ &= -\nabla_{\theta} L(z_{\text{test}}, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_{\theta} L(z, \hat{\theta})\end{aligned}$$

Influence functions

Perturbing a training input

- If we change $z = (x, y)$ to $z_\delta = (x + \delta, y)$, what will test loss change?
- $z = (x, y)$ to $z_\delta = (x + \delta, y)$ equals to delete z then add z_δ .
- New parameter:

$$\hat{\theta}_{\epsilon, z_\delta, -z} \stackrel{\text{def}}{=} \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n L(z_i, \theta) + \epsilon L(z_\delta, \theta) - \epsilon L(z, \theta)$$

- We have:

$$\begin{aligned} \left. \frac{d\hat{\theta}_{\epsilon, z_\delta, -z}}{d\epsilon} \right|_{\epsilon=0} &= \mathcal{I}_{\text{up, params}}(z_\delta) - \mathcal{I}_{\text{up, params}}(z) \\ &= -H_{\hat{\theta}}^{-1} \left(\nabla_{\theta} L(z_\delta, \hat{\theta}) - \nabla_{\theta} L(z, \hat{\theta}) \right) \end{aligned}$$

Influence functions

Perturbing a training input

- Linear approximation:

$$\hat{\theta}_{z_\delta, -z} - \hat{\theta} \approx \frac{1}{n} (\mathcal{I}_{\text{up, params}}(z_\delta) - \mathcal{I}_{\text{up, params}}(z))$$

- As $\|\delta\| \rightarrow 0$, $\nabla_{\theta} L(z_\delta, \hat{\theta}) - \nabla_{\theta} L(z, \hat{\theta}) \approx [\nabla_x \nabla_{\theta} L(z, \hat{\theta})] \delta$
- So:

$$\left. \frac{d\hat{\theta}_{\epsilon, z_\delta, -z}}{d\epsilon} \right|_{\epsilon=0} \approx -H_{\hat{\theta}}^{-1} [\nabla_x \nabla_{\theta} L(z, \hat{\theta})] \delta$$
$$\hat{\theta}_{z_\delta, -z} - \hat{\theta} \approx -\frac{1}{n} H_{\hat{\theta}}^{-1} [\nabla_x \nabla_{\theta} L(z, \hat{\theta})] \delta$$

Influence functions

Perturbing a training input

Influence of perturbing $z \rightarrow z_\delta$ on the loss at the test point z_{test} :

$$\begin{aligned}\mathcal{I}_{\text{pert,loss}}(z, z_{\text{test}}) &\stackrel{\text{def}}{=} \left. \nabla_\delta L(z_{\text{test}}, \hat{\theta}_{z_\delta, -z}) \right|_{\delta=0} \\ &= -\nabla_\theta L(z_{\text{test}}, \hat{\theta})^\top H_{\hat{\theta}}^{-1} \nabla_x \nabla_\theta L(z, \hat{\theta})\end{aligned}$$

Scaling up

Challenges

- Challenges:
 - How to calculating Inverse Hessian Matrix?
 - How to calculating influence function $\mathcal{I}_{\text{up,loss}}(z_i, z_{\text{test}})$ on all training points?
- What happens when our assumptions are violated?

Scaling up

Calculating Inverse Hessian Matrix


We have: n training points, p parameters.

- Inverse Hessian Matrix requires $O(np^2 + p^3)$
- Use Conjugate gradients (CG) requires $O(np)$
- Use Stochastic estimation
 - Key idea: Don't explicitly form H_θ^{-1} . Instead, compute $H_\theta^{-1}v$.
 - $s_{\text{test}} = H_\theta^{-1} \nabla_\theta L(z_{\text{test}}, \hat{\theta})$. Then, $\mathcal{I}_{\text{up,loss}}(z, z_{\text{test}}) = -s_{\text{test}} \nabla_\theta L(z, \hat{\theta})$

Scaling up

Influence function vs leave-one-out retraining

- Compared $-\frac{1}{n}\mathcal{I}_{\text{up,loss}}(z, z_{\text{test}})$ with $L(z_{\text{test}}, \hat{\theta}_{-z}) - L(z_{\text{test}}, \hat{\theta})$ using logistic regression, CNN model on 10-class MNIST.



Scaling up

Non-convexity

- We took $\hat{\theta}$ as the global minimum. If we have $\tilde{\theta}$ on non-convex objectives, $\tilde{\theta} \neq \hat{\theta}$. As a result, $H_{\tilde{\theta}}$ could have negative eigenvalues.
- Calculate $\mathcal{I}_{\text{up,loss}}$ using \tilde{L} :

$$\tilde{L}(z, \theta) = L(z, \tilde{\theta}) + \nabla L(z, \tilde{\theta})^T (\theta - \tilde{\theta}) + \frac{1}{2} (\theta - \tilde{\theta})^T (H_{\tilde{\theta}} + \lambda I) (\theta - \tilde{\theta})$$



Scaling up

Non-differentiable losses

- What happens when $\nabla_{\theta} L$ and $\nabla_{\theta}^2 L$, do not exists?
- Key idea: Replace origin L with smoothed version. E,g:

$$\text{Hinge}(s) = \max(0, 1 - s)$$

$$\approx \text{SmoothHinge}(s, t) := t \log \left(1 + \exp \left(\frac{1 - s}{t} \right) \right)$$



Application

- ① Understanding model behavior
- ② Adversarial training examples
- ③ Debugging domain mismatch
- ④ Fixing mislabeled examples

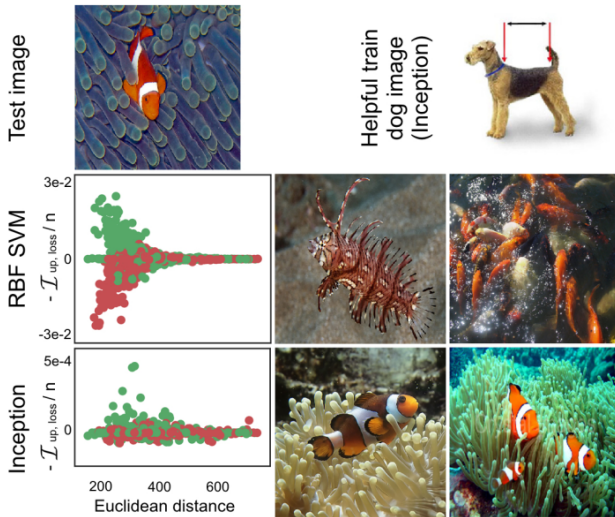
Applications

Understanding model behavior

- Model 1: Inception v2 with all but the top layer frozen (used pre-trained from Keras).
- Model 2: SVM with RBF kernel.
- Task: Binary classification of fish and dog.

Applications

Understanding model behavior



Applications

Adversarial training examples

- There exists some paper generating some adversarial test images that are visually indistinguishable but can fool a classifier.
- We demonstrate we can craft adversarial training images that can flip a model's prediction.
- Basically, the idea is iterating on training images on the direction of influence function.

Applications

Adversarial training examples

- Data same as fish and dog
 - Origin correctly classified 591/600 test images.
 - For each test image, find only one training image and do 100 iterations.
 - 335 (57%) of the testing images were flipped.
 - Also, attack on one training image can influence multiple test images.

Applications

Adversarial training examples


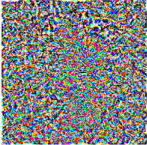




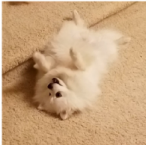

A small perturbation to one **training** example:

Label: Fish

$+ \epsilon \cdot$




Label: Fish

Can change multiple **test** predictions:

				
				
Orig (confidence): Dog (97%)	Dog (98%)	Dog (98%)	Dog (99%)	Dog (98%)
New (confidence): Fish (97%)	Fish (93%)	Fish (87%)	Fish (60%)	Fish (51%)

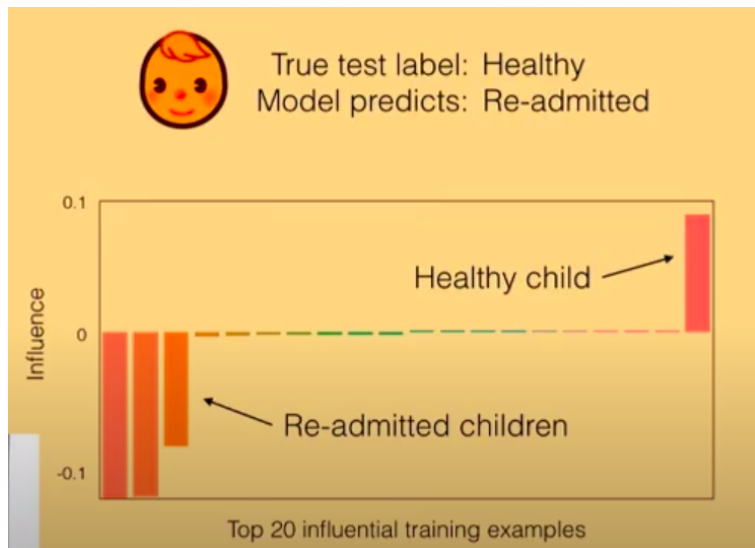
Applications

Debugging domain mismatch

		Original		Modified
Healthy + re-admitted adults		~20k	→ same →	~20k
Healthy children		21	→ -20 →	1
Re-admitted children		3	→ same →	3

Applications


Debugging domain mismatch



Applications

Fixing mislabeled examples

- Only have training set. Find example with largest loss.
- We measure the influence of z_i with $\mathcal{I}_{\text{up,loss}}(z_i, z_i)$ which approximates the error incurred on z_i if we remove z_i from training data.



Conclusion

- We can better understand the behavior of a model by looking at how it was derived from training data.
- Influence functions let us do this efficiently.
- Key: Differentiate through training and rely on asymptotics.
- Locality allows us to get a closed form expression. Can we get at a global notion of influence?
- Much more work to be done on these black box model tool.

Thank you for your attention!

